



數系天地
勤篤求真

Geometric Learning

Academy of Mathematics and Systems Sciences, CAS

May 18, 2023

Contents

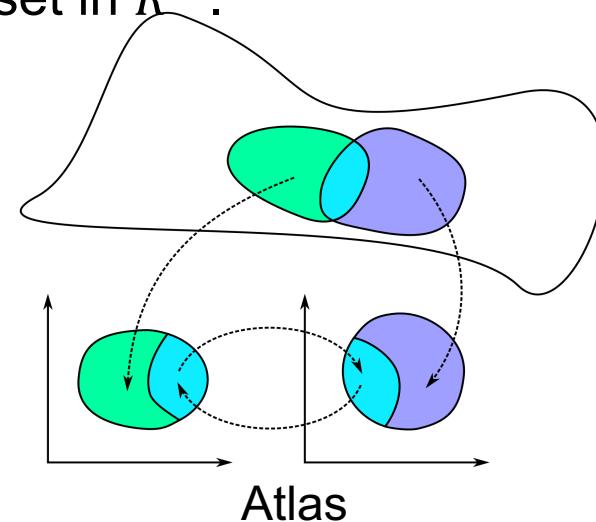
- Introduction
 - Geometry and Geometric Data: a Brief Introduction
 - Manifolds and Potential Structures
 - Point Clouds, Meshes and High-dimensional Geometric Data
 - Machine learning for Object Matching and Manifold Learning
 - Gromov-Hausdorff Metric
 - Functional Map
 - Isometry-based Methods
- AI4Geo

Differentiable Manifold

Differentiable Manifold

M is a differentiable manifold with dimension m if:

- M is a Hausdorff and second countable topological space.
- M has a differentiable atlas $\{\phi_i\}$. Image of ϕ_i is an open set in \mathbb{R}^m .



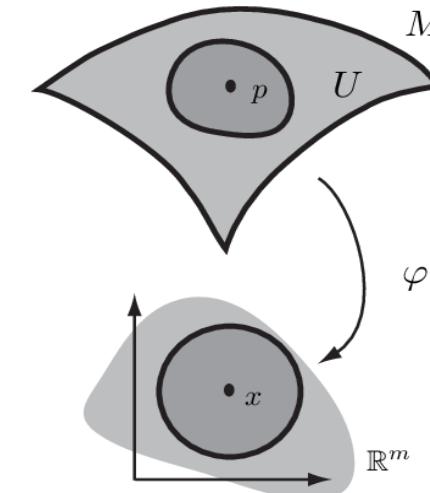
Embedding

$$n \leq 2m + 1$$

Submanifold of Euclidian Space

M is a submanifold of \mathbb{R}^n with dimension m if:

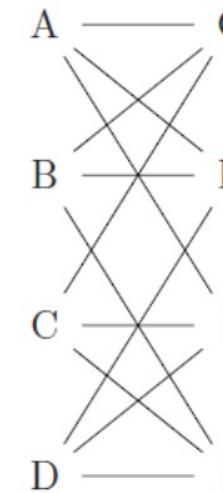
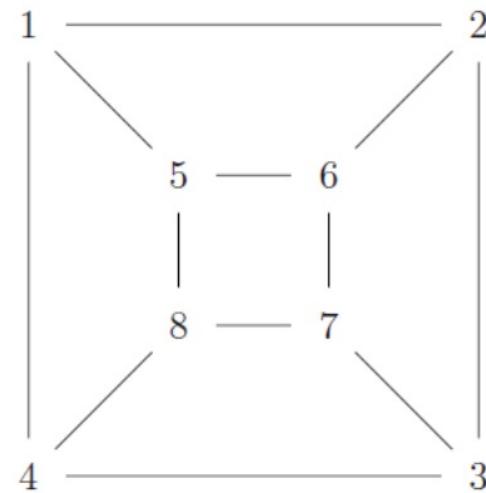
- M is a subset of \mathbb{R}^n
- $\forall x \in M, \exists x \in B, s.t. B \cap M$ and U and diffeomorphic, where B is a neighbor of $x \in \mathbb{R}^n$ and U is an open set in \mathbb{R}^m .



Potential Structures on Manifold

- Metric space and isomorphism.
 - (M, d) is a set M equipped with distance function d .
 - $(M_1, d_1), (M_2, d_2)$ are isomorphic if there exists an invertible map $T: M_1 \rightarrow M_2$ satisfy

$$d_1(x, y) = d_2(T(x), T(y))$$



$f(A) =$	1
$f(B) =$	6
$f(C) =$	3
$f(D) =$	8
$f(G) =$	2
$f(H) =$	5
$f(I) =$	4
$f(J) =$	7

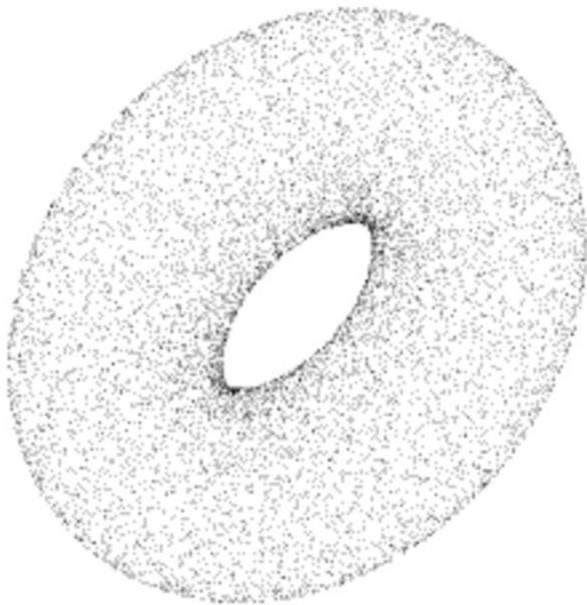
A manifold M can be a metric space.
For example, Riemannian manifold.

Potential Structures on Manifold

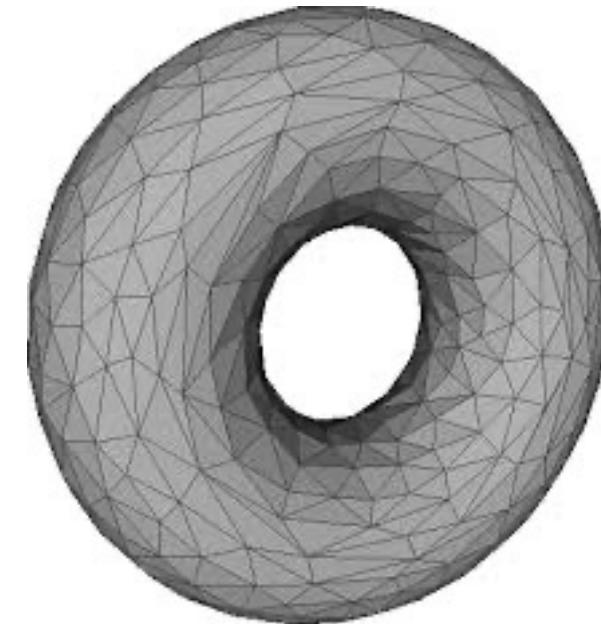
- A measure space is a triple (X, \mathfrak{F}, μ) , where
 - X is a set
 - \mathfrak{F} is a σ -algebra on X
 - μ is a measure on (X, \mathfrak{F})
- A metric measure space is a triple (X, d_X, μ_X) where
 - (X, d_X) is a compact metric space
 - μ_X is a Borel probability measure on X and μ_X has full support:
$$supp[\mu_X] = X$$

A manifold M can be a measure space,
even a metric measure space.

3D Geometric Data



Point clouds: points only



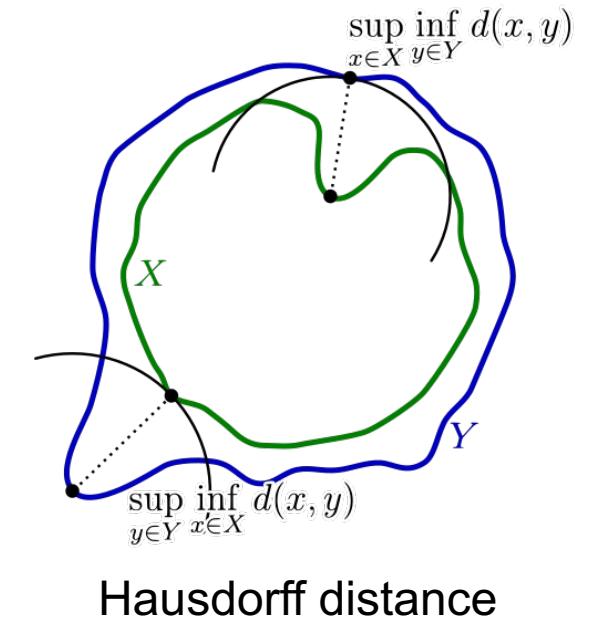
Mesh: points + edges + faces

How to recognize an object from geometric data?
How to match the shapes of objects?

Object Matching

- Comparing objects by **Hausdorff distance**:
 - Assume A, B are compact subsets of metric space (Z, d) . The Hausdorff distance $d_{\mathcal{H}}$ between A, B is defined by

$$d_{\mathcal{H}}^Z(A, B) := \max\left(\sup_{a \in A} \inf_{b \in B} d(a, b), \sup_{b \in B} \inf_{a \in A} d(a, b)\right)$$



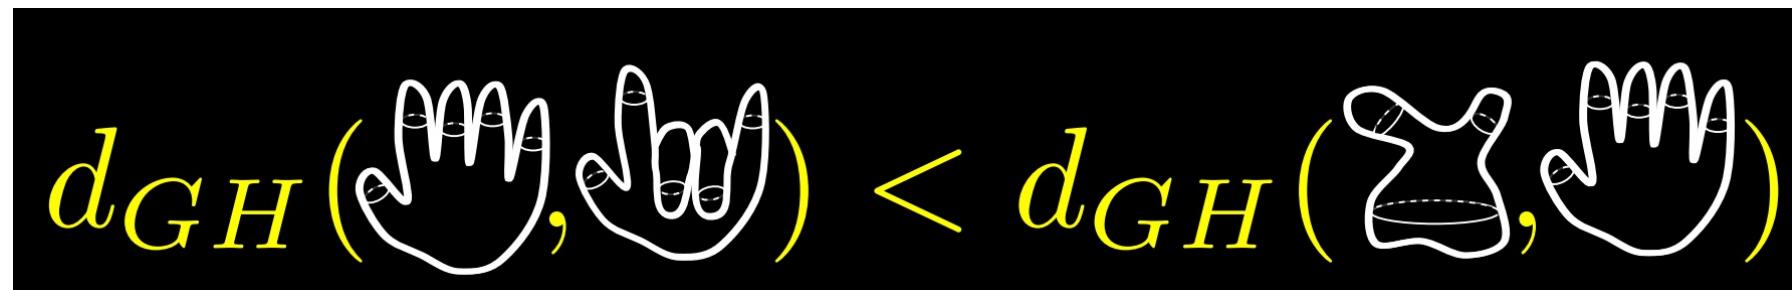
- Comparing objects by **intrinsic distance**:
 - Assume $(A = \{a_i\}, d_A), (B = \{b_j\}, d_B)$ are metric space. Then d_A and d_B correspond to a matrix.
 - Measure the difference by measuring the difference of these matrices.

Object Matching

- The **Gromov-Hausdorff distance** combines these ideas. Suppose X, Y are two compact metric spaces, then the Gromov-Hausdorff distance is defined as:

$$d_{\mathcal{GH}}(X, Y) := \inf_{Z, f, g} d_{\mathcal{H}}^Z(f(X), g(Y))$$

where f, g are isometric embeddings into the metric space Z .


$$d_{\mathcal{GH}}(\text{Hand}, \text{Hand}) < d_{\mathcal{GH}}(\text{Hand}, \text{Torus})$$

- Surprisingly, $d_{\mathcal{GH}}(X, Y)$ is a metric of all compact metric spaces quotient isometry equivalence.

Object Matching

- A correspondence R of X and Y is a subset of $X \times Y$ s.t.

- $\forall x \in X, \exists y \in Y$ s.t. $(x, y) \in R$
- $\forall y \in Y \exists x \in X$ s.t. $(x, y) \in R$

- The following equation gives the way to calculate $d_{\mathcal{GH}}$:

$$d_{\mathcal{GH}}(X, Y) = \frac{1}{2} \inf_{R \in \mathcal{R}(X, Y)} \sup_{\substack{x_1, x_2 \in X \\ y_1, y_2 \in Y \\ s.t. (x_i, y_i) \in R}} \Gamma_{X, Y}(x_1, y_1, x_2, y_2)$$

where $\mathcal{R}(X, Y)$ is the collection of correspondence between X and Y ,
 $\Gamma_{X, Y}(x, y, x', y') = |d_X(x, x') - d_Y(y, y')|$.

Object Matching

- However, finding the best correspondence is **NP-hard!**
- There are several ways to solve the problem:
 - Solve landmark correspondence
 - MDS+PCA [Bronstein et al. (2006)]
 - Functional maps [Sun et al. (2009)]

Functional Maps [Ovsjanikov et al. (2012)]

- A mapping $f: M \rightarrow N$ induces a linear mapping between function space:
$$f_{\#}: \{g: N \rightarrow R\} \rightarrow \{h: M \rightarrow R\}$$
$$g \rightarrow g \circ f$$
- Choose a finite-dimensional space of $L_2(M)$ and $L_2(N)$ and a basis in them, the $f_{\#}$ corresponds to a matrix C .
- Consider eigen functions of Laplacian over a compact Riemannian manifold M :

$$\Delta_M u = \lambda u$$

If we choose the basis as the first k eigen functions of Δ_M and Δ_N (denote the i -th eigen function as ϕ_i, ψ_i), then we can recover the most likely mapping f from its matrix C .

Functional Maps [Sun et al. (2009)]

- These functions are fingerprints of this Riemannian manifold. To be specific,
 - $T: M \rightarrow N$ is isomorphic between Riemannian manifold M, N
 $\Leftrightarrow \{\text{eigen values of } \Delta_M\} = \{\text{eigen values of } \Delta_N\}$ and $\phi_i = \psi_i \circ T$
 - By heat equation, we have that:

$$\delta_x(\cdot) = \lim_{t \rightarrow 0} \sum^{\infty} e^{-\lambda t} \phi_i(x) \phi_i(\cdot)$$

- As a result, $\langle f_{\#}(\delta_{(\cdot)}), \delta_{(\cdot)} \rangle \approx |\Psi \mathcal{C} \Phi^T|$, $|\Psi \mathcal{C} \Phi^T|^\wedge$ gives the probability of image of $f(x)$, where

$$\Phi = [\phi_i(x_j)]_{ji}, \quad \Psi = [\psi_i(y_j)]_{ji}$$

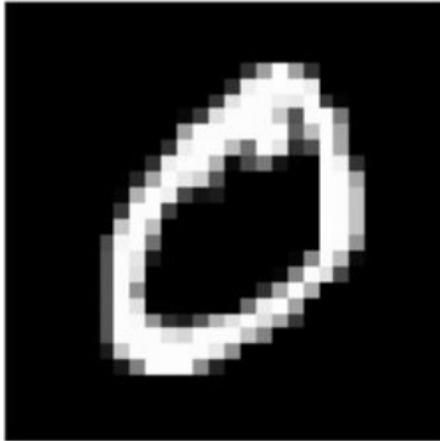
and “ $^\wedge$ ” means column-wise normalization, $\{x_j\}$ is the data set observed on the manifold.

High-dimensional Geometric Data

High-dimensional data

lie along

Low-dimensional manifolds



Dimension: 28×28

3	4	2	1	9	5	6	2	1	8
8	9	1	2	5	0	0	6	6	4
6	7	0	1	6	3	6	3	7	0
3	7	7	9	4	6	6	1	8	2
2	9	3	4	3	9	8	7	2	5
1	5	9	8	3	6	5	7	2	3
9	3	1	9	1	5	8	0	8	4
5	6	2	6	8	5	8	8	9	9
3	7	7	0	9	4	8	5	4	3
7	9	6	4	7	0	6	9	2	3

Dimension: 2

High-dimensional Geometric Data

High-dimensional data



lie along

Low-dimensional manifolds



Dimension: $178 \times 218 \times 3$

Dimension: 100

How to lower the dimension of data in the meanwhile
preserve the structure of manifolds?

Manifold Learning

- From high-dimensional data to lower-dimensional representation:

$$\{x_i \in R^n\} \rightarrow \{y_i \in R^m\}, \ m \ll n$$

where x_i corresponds to y_i . In the meanwhile, the **structure** of the manifold should be preserved.

- MDS: preserve extrinsic metric

$$\min \sum_{i,j} \|d_{R^n}(x_i, x_j) - d_{R^m}(y_i, y_j)\|_2^2$$

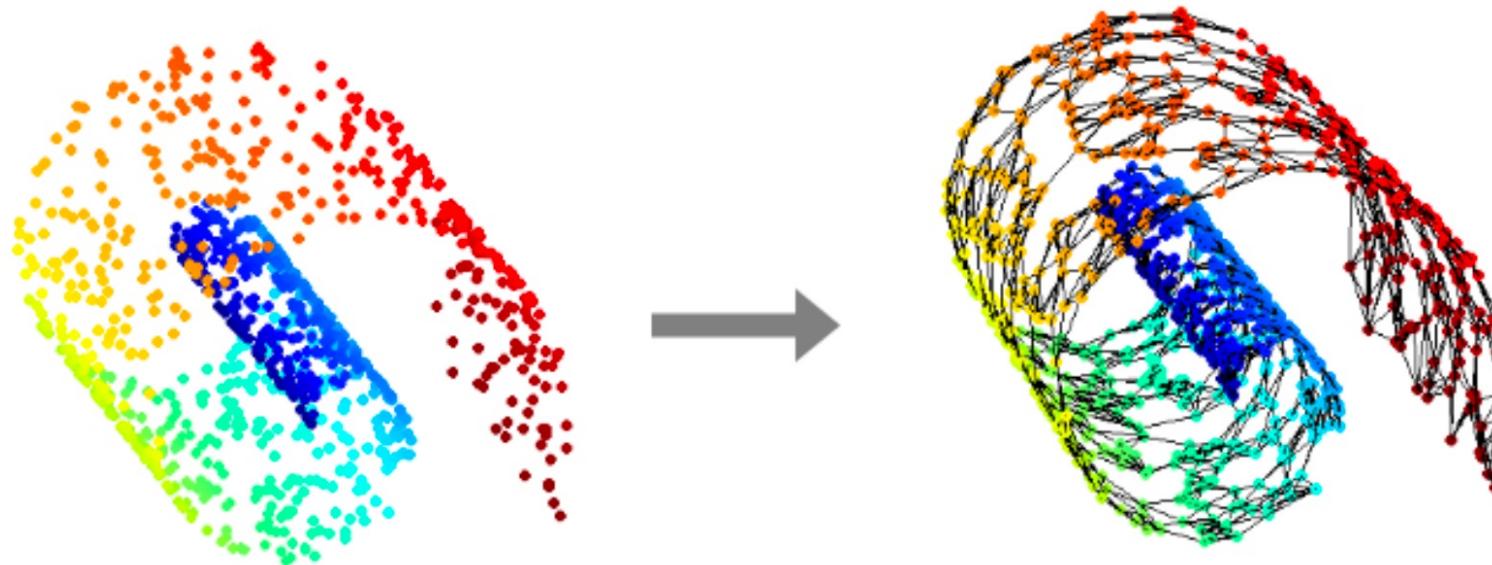
- Isomap: preserve intrinsic metric [Tenenbaum et al. (2000)]

$$\min \sum_{i,j} \|d_M(x_i, x_j) - d_{R^m}(y_i, y_j)\|_2^2$$

Manifold Learning

- Isomap: preserve intrinsic metric

$$\min \sum_{i,j} \|d_M(x_i, x_j) - d_{R^m}(y_i, y_j)\|_2^2$$



Construct graph structure for geodesic length

Contents

- Deep Learning for 3D Shapes
 - Learning representation of points clouds
 - Recognition
 - Segmentation
 - Learning developable approximation
 - Learning for shape matching
- Deep Learning for High-dimensional Geometry
 - Wasserstein Autoencoders
 - Gromov-Wasserstein Autoencoders

Learning Representation of Point Clouds

- Distinguishing different objects by traditional methods needs an estimation of metric, which makes it **time-consuming**.
- Traditional segmentation methods have been developed for certain usages. **No general pipeline** for segmentation.
- Now that there are many available **datasets of point clouds and their classification** (such as “chairs”, “cups” and so on) and semantic (or other standards) segmentation (e.g. chair leg, cup handle).
- Deep learning can be used to **recognize objects and segmentation**.

Learning Representation of Point Clouds

- **Loss function for 3D object classification:** the softmax classification loss + a regularizer.
- Let's denote the dataset by $\{(x_i, y_i)\}$, where x_i is a point cloud and y_i is the label.

$$L = \sum_i \log \frac{e^{y_i}}{\sum_j e^{f_\theta(x)_j}}$$

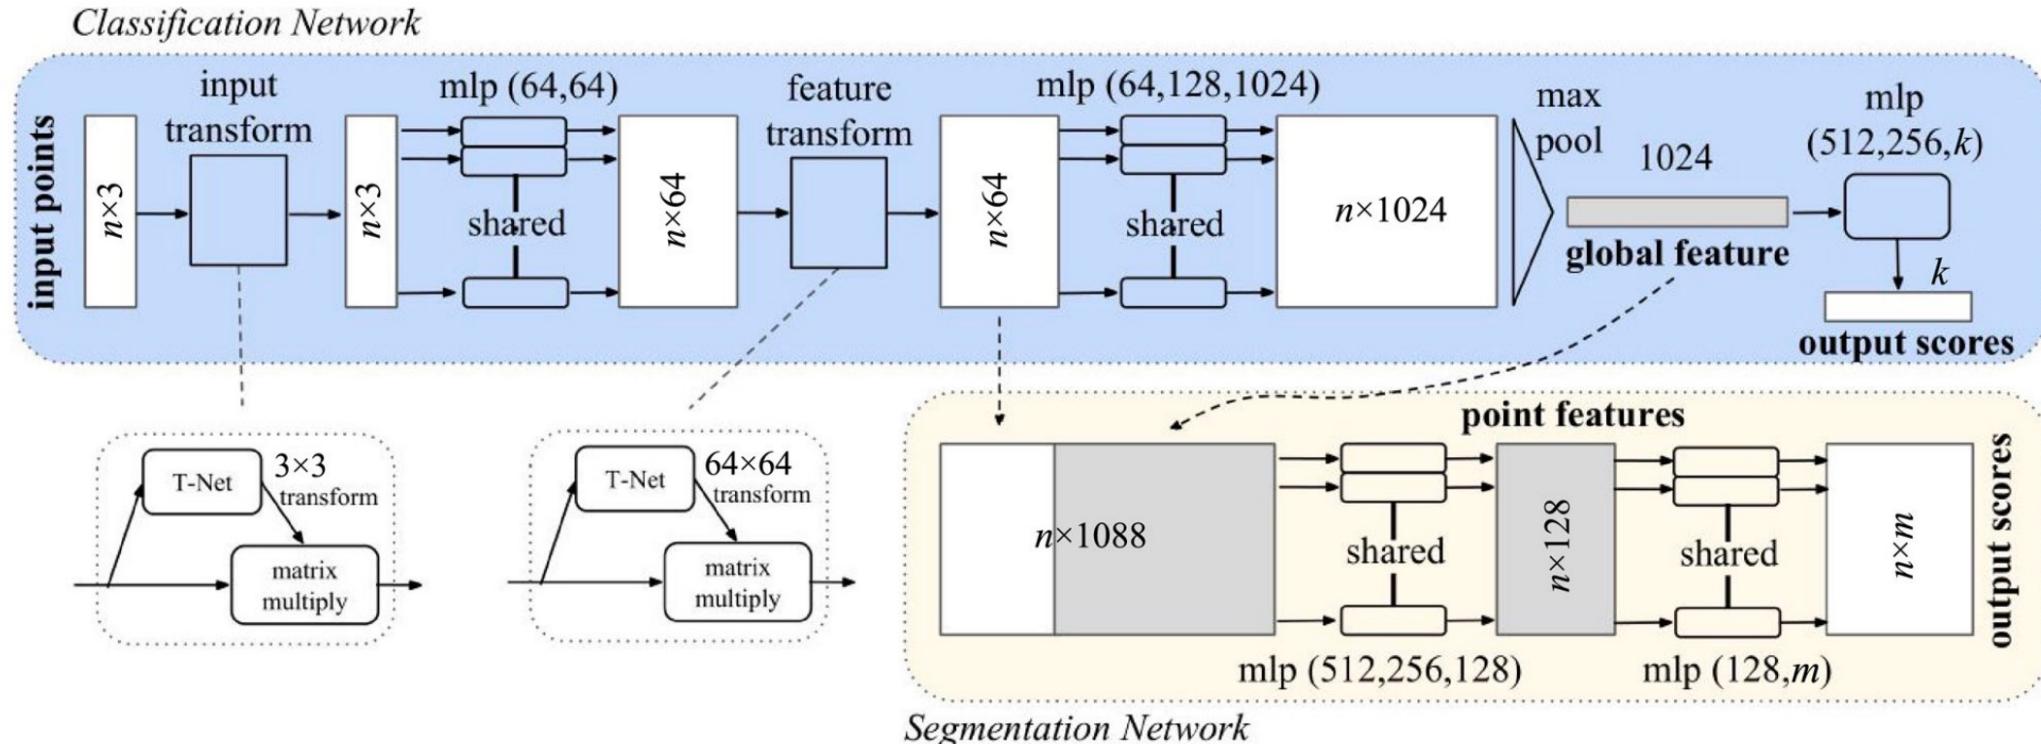
- **Loss function for 3D object segmentation:** the Mean Intersection over Union (MIoU) loss.
- Let's denote the probability of a point of class i be classified into class j by $p_{ij}(\theta)$. Then the MIoU loss is defined as:

$$\text{MIoU} = \frac{1}{k} \sum_{i=1}^k \frac{p_{ii}(\theta)}{\sum_j p_{ij}(\theta) + \sum_j p_{ji}(\theta) + p_{ii}(\theta)}$$

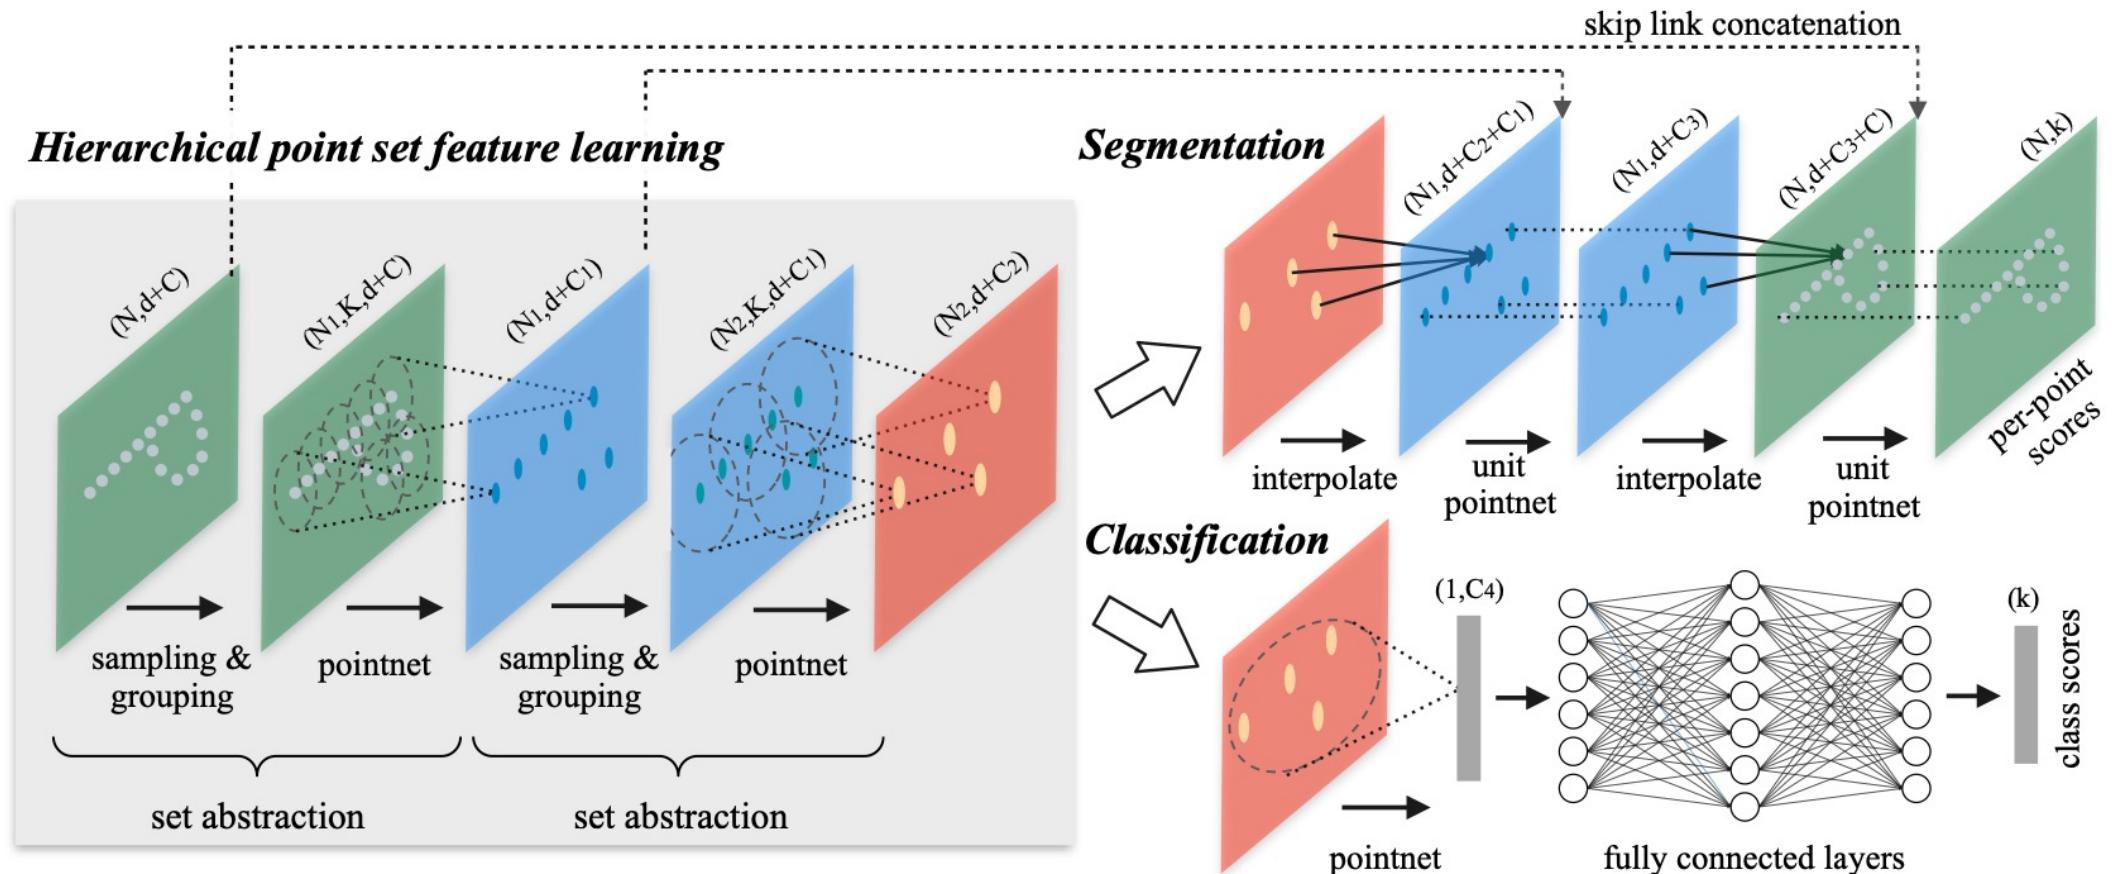
which can be interpreted as the mean of
$$\frac{\text{True Positive}}{\text{False Negativ} + \text{False Positive} + \text{True Positive}}$$

Learning Representation of Point Clouds: PointNet

- PointNet: learn **global features** and the **point features** of the point cloud.
- The global feature can be used to do **3D object classification** while the point features can be used to do **segmentation** [Qi et al. (2017)].

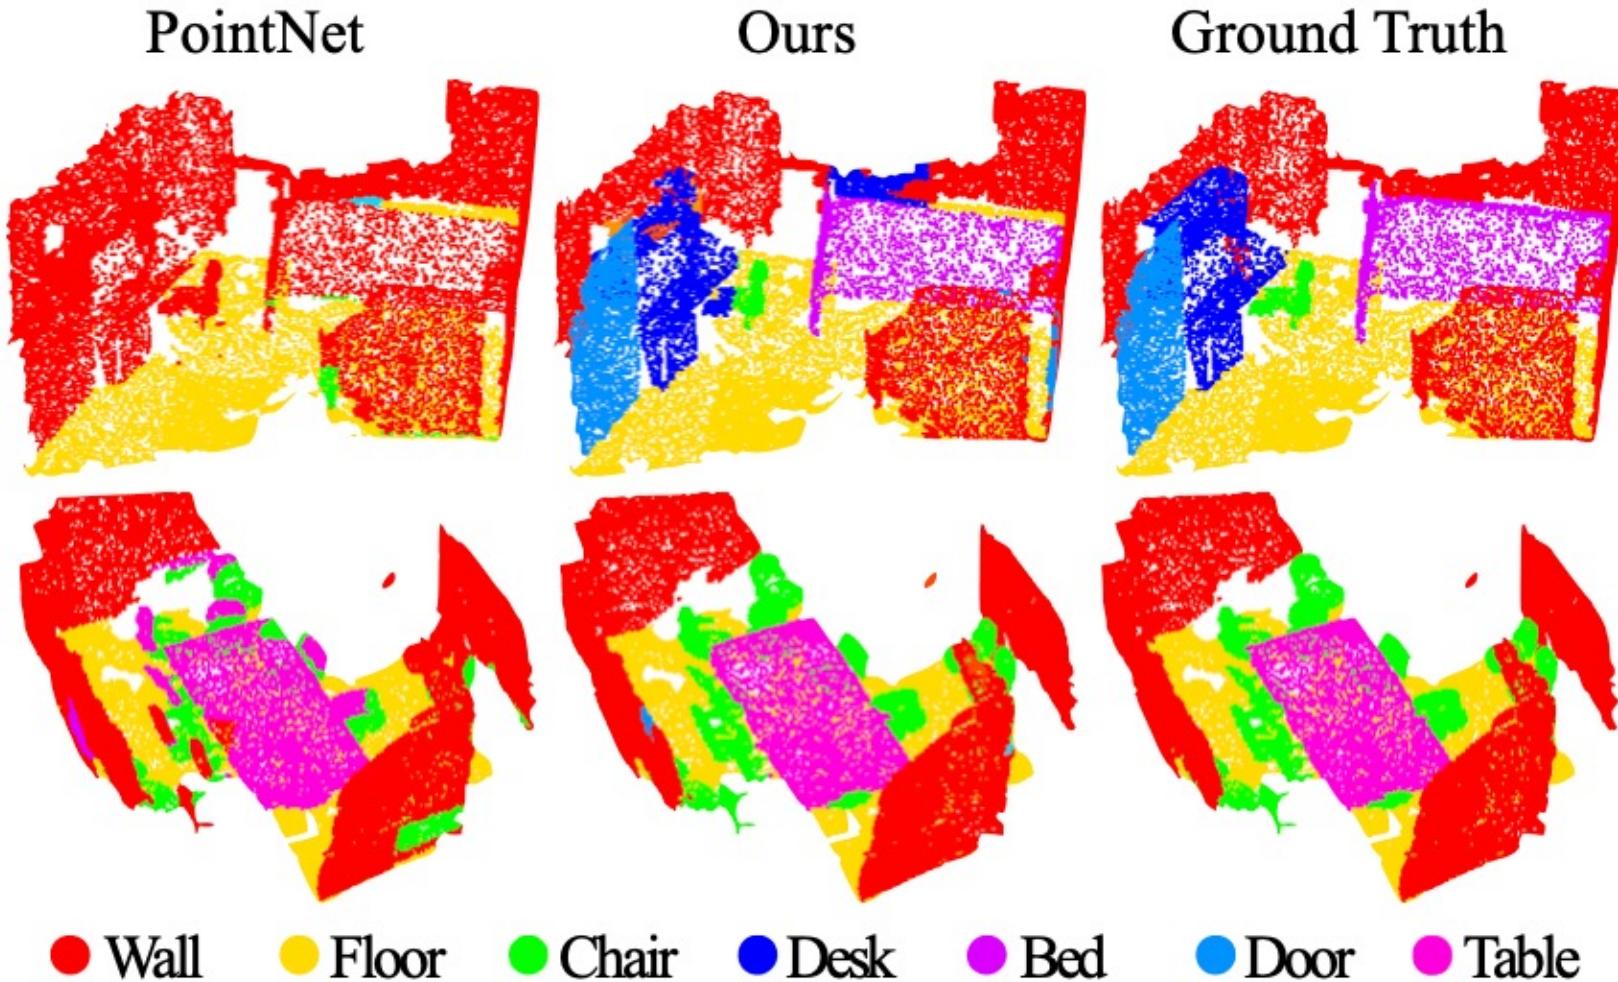


Learning Representation of Point Clouds: PointNet++



PointNet++ adopts a hierarchical structure to further capture local features [Qi et al. (2017)]

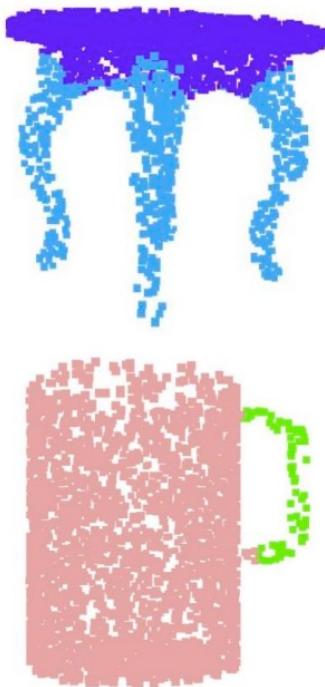
PointNet vs PointNet++



PointNet vs PointNet++

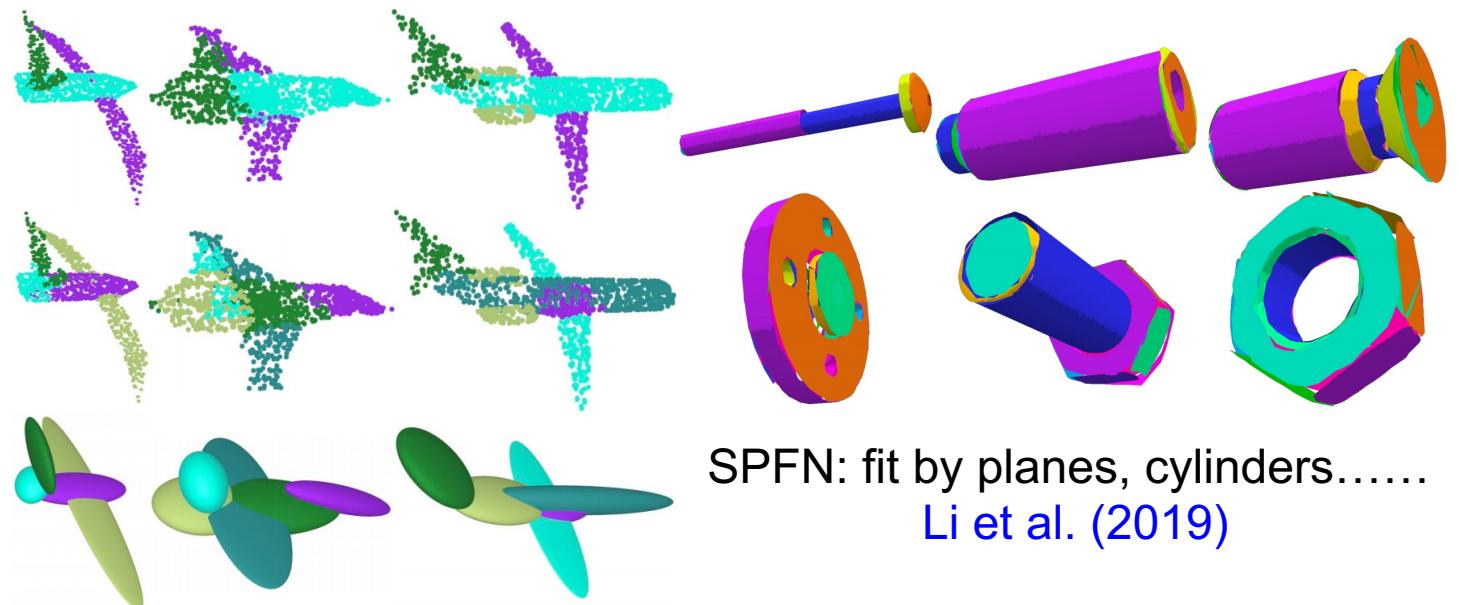
Other Segmentation

Fitting by semantics



PointNet

Fitting by primitives

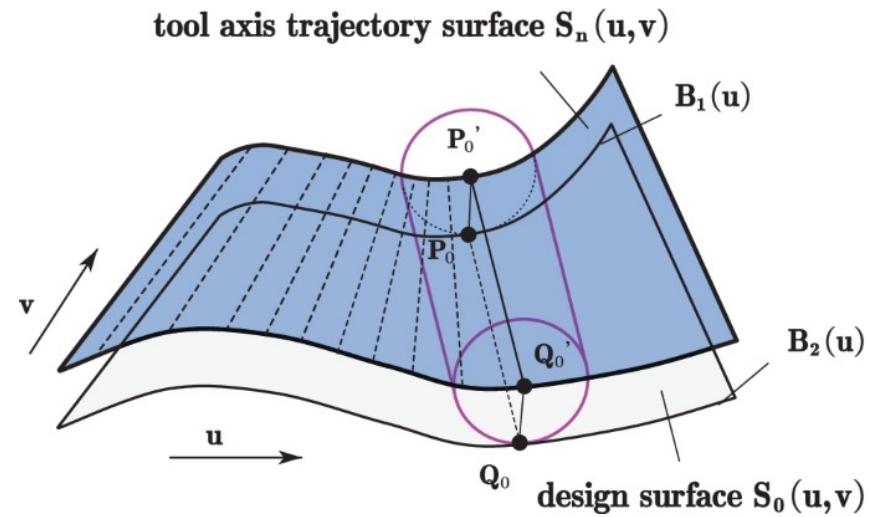
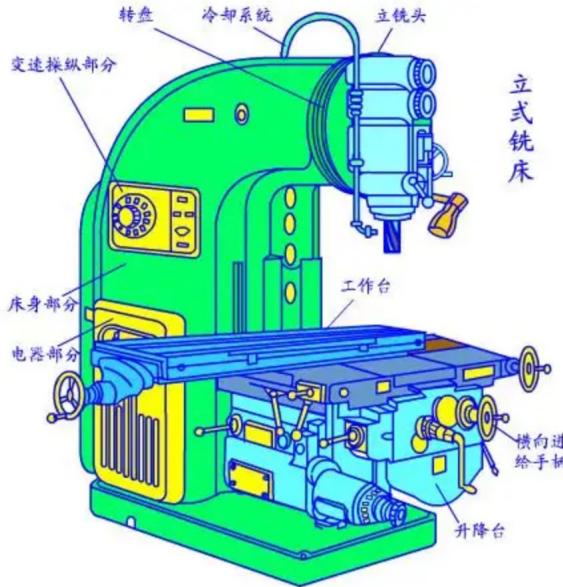


SPFN: fit by planes, cylinders.....
Li et al. (2019)

PRIFIT: fit by ellipsoids
Sharma et al. (2022)

Different datasets, additional loss terms for specific tasks,
and similar structure (based on PointNet++)

Learning Developable Approximation



数控铣床

铣刀切割

可展曲面

铣刀切割 \Leftrightarrow 可展曲面

Developable Approximation

- **Developable approximation:** transform a complex surface to a set of quasi-developable surfaces.



Mitani & Suzuki

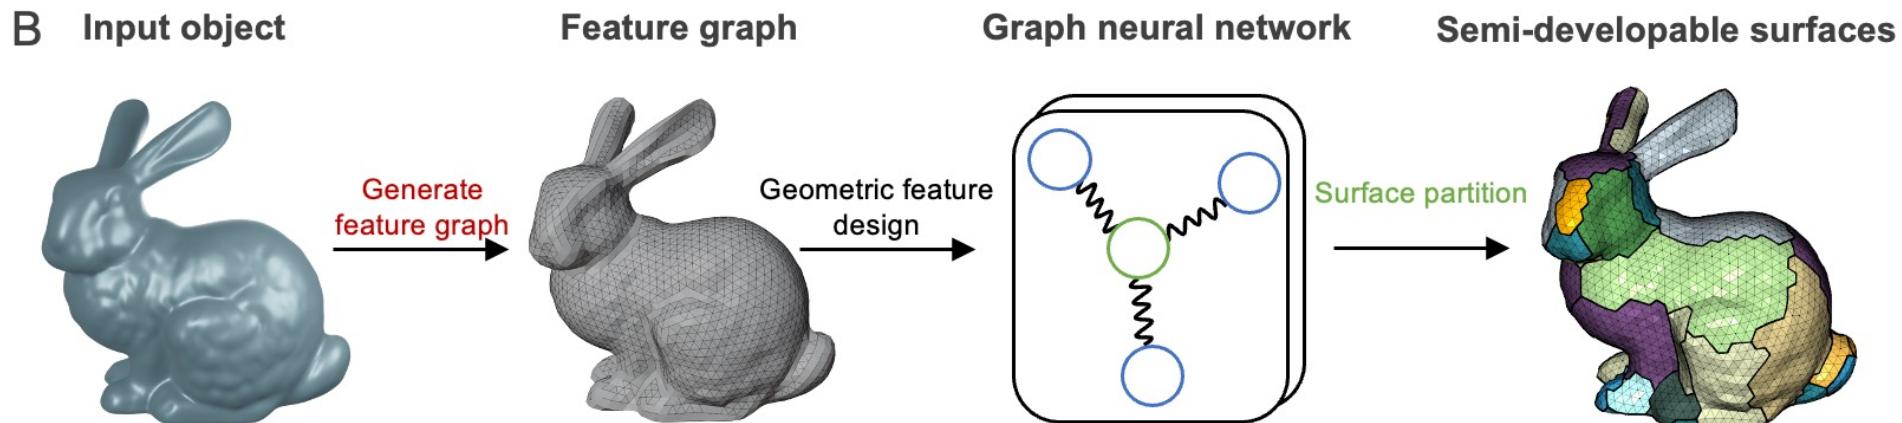
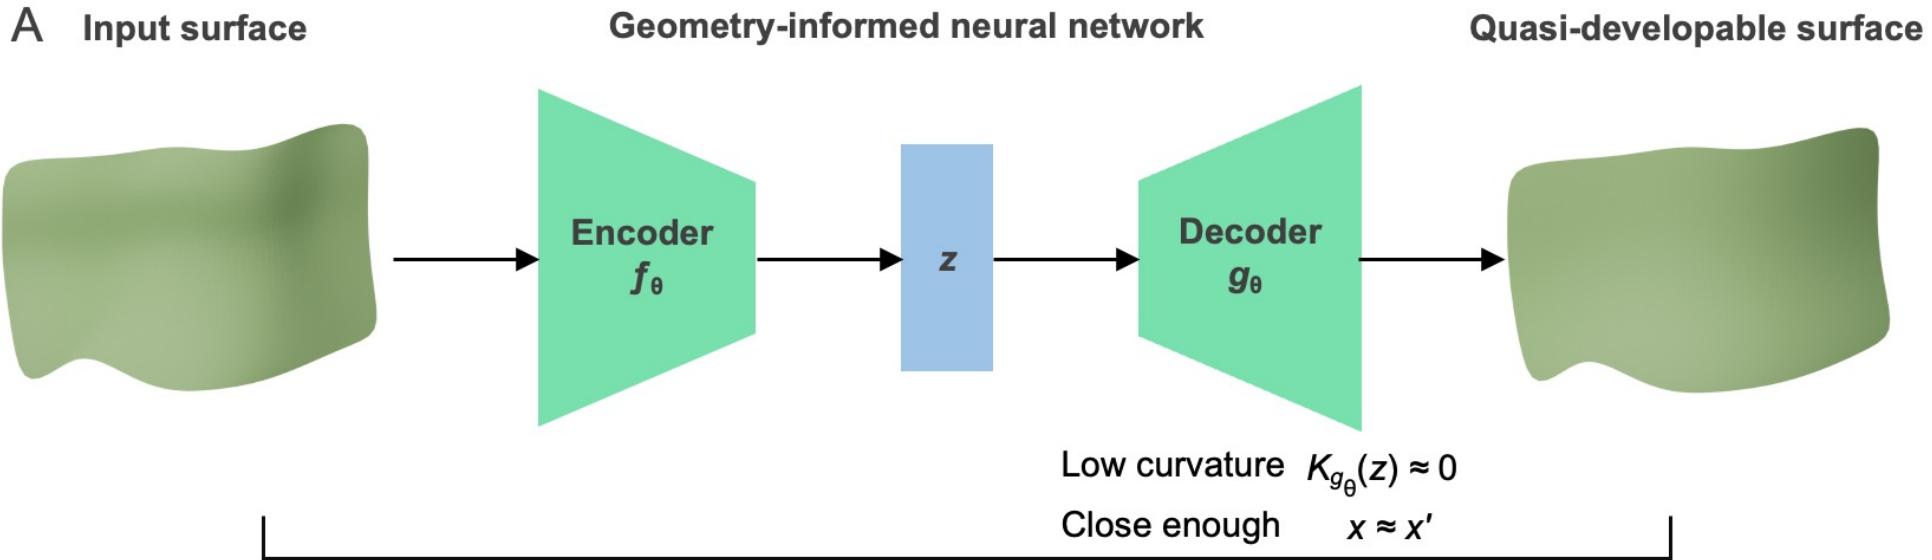
Shatz et. al.

Stein et. al.

Binninger et. al.

Zhao & Fang

Learning Developable Approximation



Segmentation



Segmentation + Lower curvature

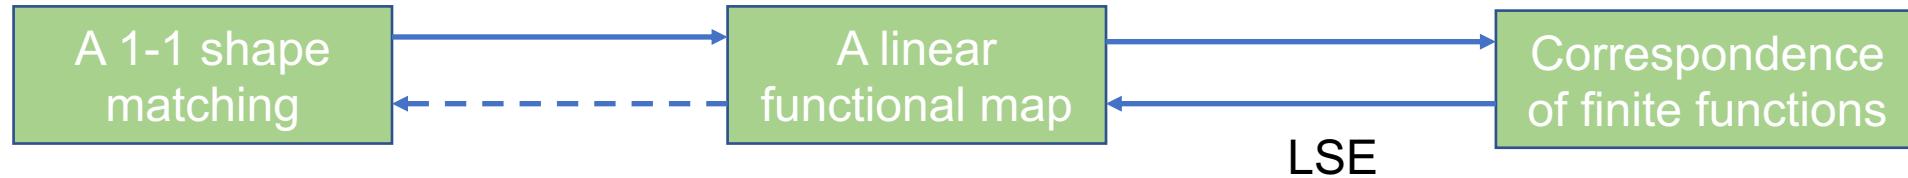


Learning for Shape Matching

- Recall that $d_{\mathcal{GH}}(X, Y) = \frac{1}{2} \inf_{R \in \mathcal{R}(X, Y)} \sup_{\substack{x_1, x_2 \in X \\ y_1, y_2 \in Y \\ s.t. (x_i, y_i) \in R}} \Gamma_{X, Y}(x_1, y_1, x_2, y_2)$. To calculate the $d_{\mathcal{GH}}$ is to find the **best correspondence**, which is NP-hard.
- Note that we have a lot of data of the best correspondence.
- It is natural to use **supervised learning** to learn the correspondence then accelerate the solving process [Boscaini et al. (2016)].

Learning for Shape Matching

- **Functional map**: a linear map $L: \{g: N \rightarrow R\} \rightarrow \{h: M \rightarrow R\}$ can give a way of shape matching, which can be estimated by **LSE** if we have already known that $L(f_i) = g_i$.



- Consequently, learning a **shape matching correspondence** is to learn a **functional correspondence**:

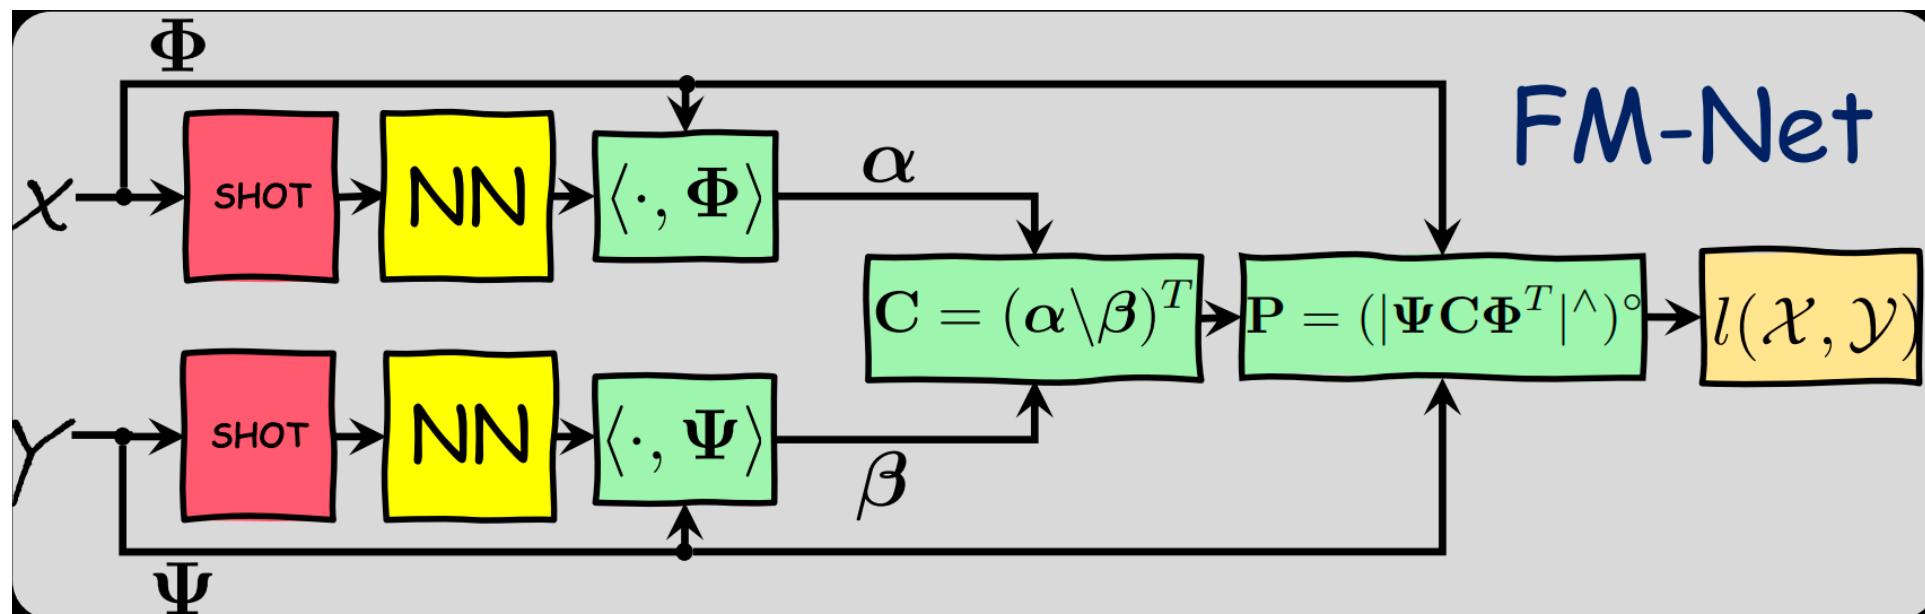
$$L(f_{i,\theta}) = g_{i,\theta}$$

- Luckily, NN has strong ability to represent functions.

Learning for Shape Matching [Litany et al. (2017)]

- Dataset: $\{(X, Y, \pi^*)\}$ where X, Y are meshes and $\forall x \in V(X), \pi^*(x)$ gives the best correspondent point of x .
- Loss function:

$$l_{sup}(X, Y) = \sum_{i \in V(X)} \sum_{j \in V(Y)} p_{ij,\theta} d_y^2(j, \pi^*(i))$$



Learning for Shape Matching [Halimi et al. (2019)]

- Dataset: $\{(X, Y, \pi^*)\}$ where X, Y are meshes and $\forall x \in V(X), \pi^*(x)$ gives the best correspondent point of x .
- Loss function:

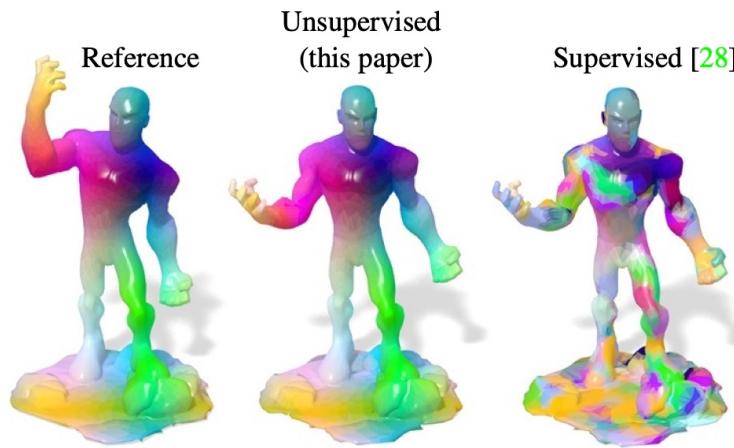
$$l_{uns}(X, Y) = \|D_X - PD_Y P^T\|$$

where D_X, D_Y are distance matrices of X and Y .

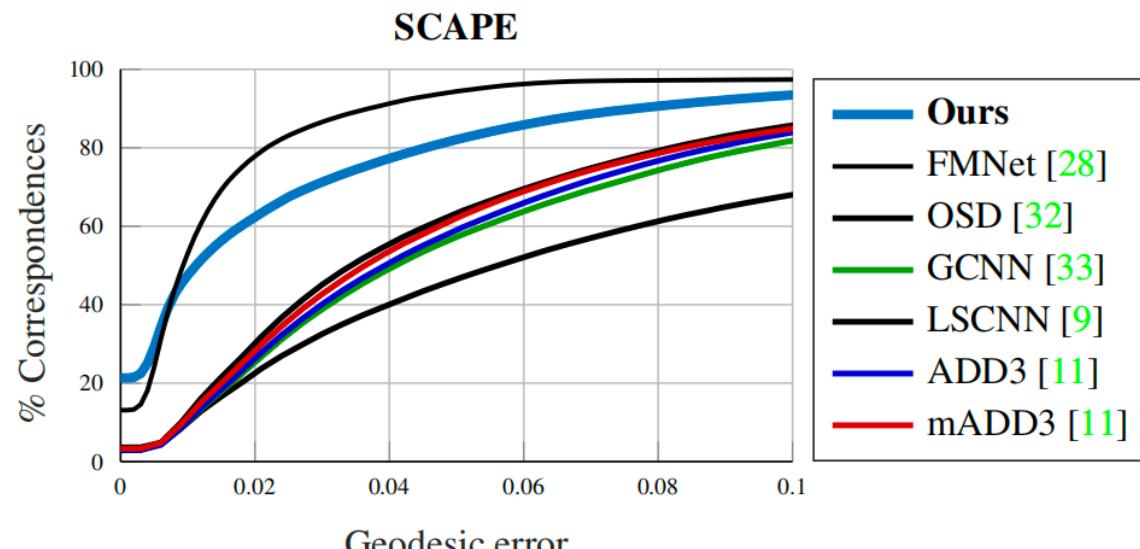
Learning for Shape Matching



Results of FMNet



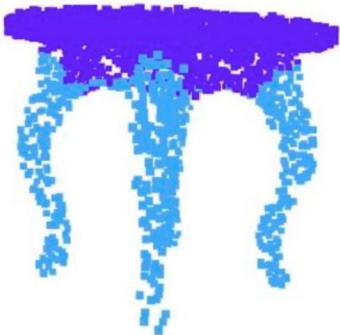
Unsupervised vs Supervised on single



Unsupervised vs Supervised

Wasserstein Metric and Gromov-Wasserstein metric

Data of different type



Evenly distributed,
low-dimensional data



Unevenly distributed,
high-dimensional data

Modeling and metric

- Metric space
 - Hausdorff distance
 - Gromov-Hausdorff metric

- Manifold assumption
- Metric measure space
 - Wasserstein metric
 - Gromov-Wasserstein metric

Wasserstein Metric

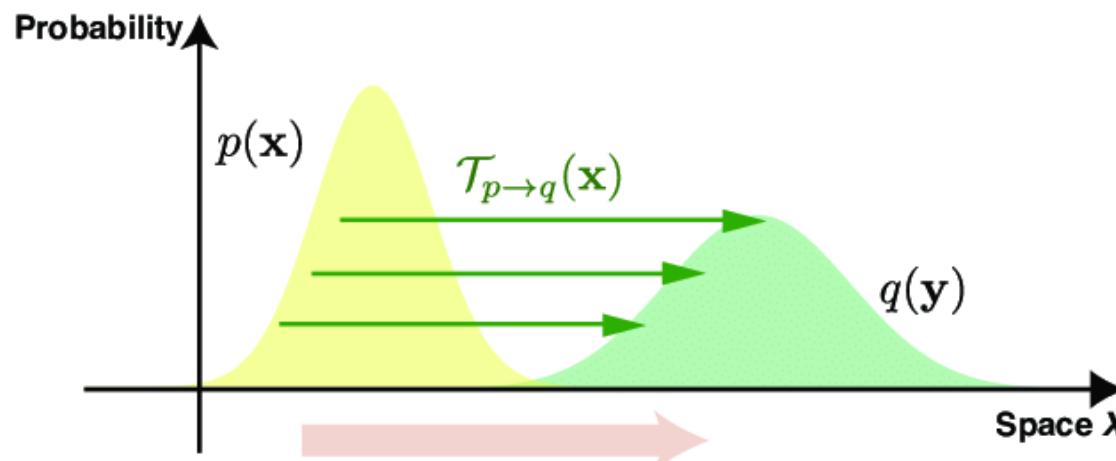
- Given two probability measure μ_A, μ_B with compact support A_0, B_0 in R^n , the L^k Wasserstein distance is defined as

$$\mathcal{W}_k(\mu_A, \mu_B)^k := \inf_{\mu \in \mathcal{M}(\mu_A, \mu_B)} \int \|x - y\|^k d\mu(x, y)$$

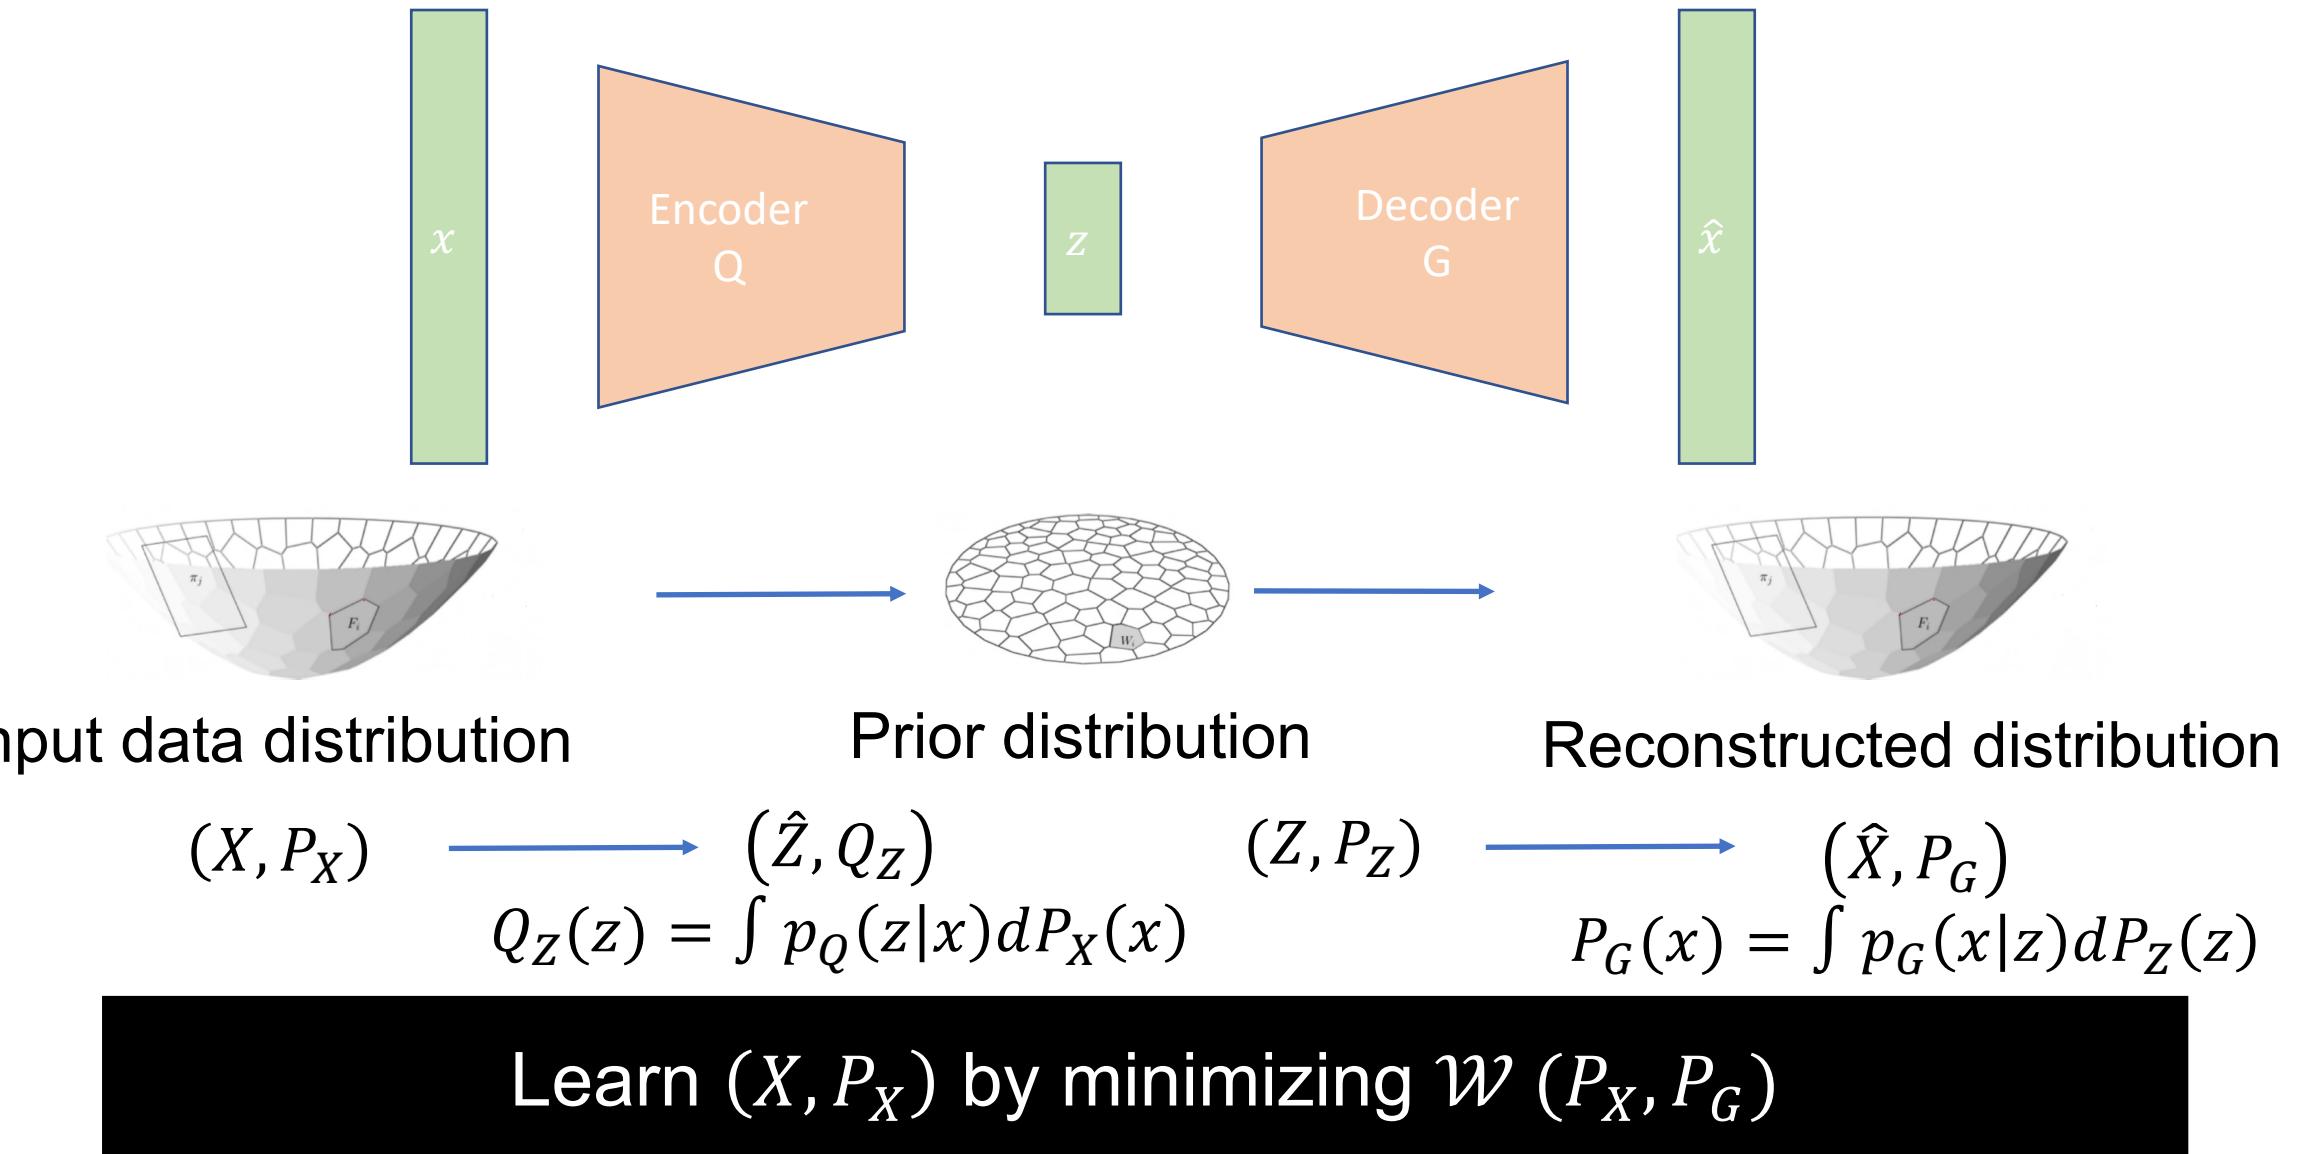
$\mathcal{M}(\mu_A, \mu_B)$ is the set of coupling of μ_A and μ_B :

$$\mathcal{M}(\mu_A, \mu_B) := \{\text{measure } \mu \text{ on } A_0 \times B_0, \text{ s.t. } \mu(A_0 \times B) = \mu_B(B), \mu(A \times B_0) = \mu_A(A)\}$$

- Wasserstein metric is also known as **mass transportation distance**.



Wasserstein Autoencoder [Tolstikhin et al. (2017)]



Wasserstein Autoencoder

For P_G as defined above with deterministic G and any function $G: Z \rightarrow X$

$$\inf_{\gamma \in \mathcal{M}(P_X, P_G)} E_{\gamma(X, Y)}[c(x, y)] = \inf_{Q: Q_Z = P_Z} E_{P_X} E_{Q(Z|X)}[c(X, G(z))]$$

- Loss function can be written as:

$$L = E_{P_X} E_{Q(Z|X)}[c(X, G(z))] + \lambda D_Z(Q_Z, P_Z)$$

where $D_Z(\cdot, \cdot)$ is an arbitrary divergence. For example:

- JS divergence

$$\text{JS}(P, Q) = \frac{1}{2} \text{KL}\left(P, \frac{P + Q}{2}\right) + \frac{1}{2} \text{KL}\left(Q, \frac{P + Q}{2}\right)$$

- Maximum mean discrepancy (MMD)

$$\text{MMD}[\mathcal{F}, P, Q] := \sup_{f \in \mathcal{F}} (E_P[f(x)] - E_Q[f(y)])$$

When \mathcal{F} is an RKHS of a positive-definite reproducing kernel $k: Z \times Z \rightarrow R$,

$$\text{MMD}_k(P, Q) = \left\| \int_Z k(z, \cdot) dP(z) - k(z, \cdot) dQ(z) \right\|_{\mathcal{H}_k}$$

where \mathcal{H}_k is the RKHS of real-valued functions mapping Z to R .

Wasserstein Autoencoder

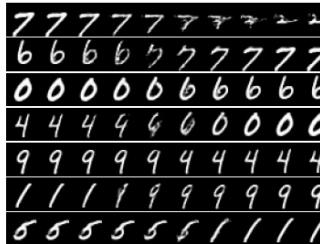
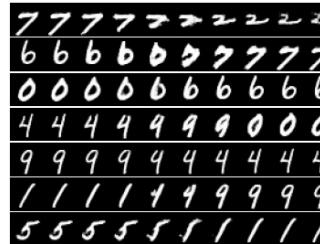
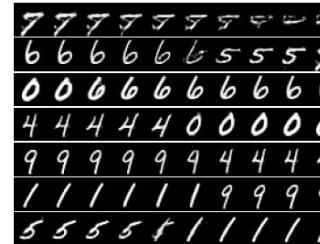
	VAE	WAE-MMD	WAE-GAN
Test interpolations			
Test reconstructions			
Random samples			

Figure 2: VAE (left column), WAE-MMD (middle column), and WAE-GAN (right column) trained on MNIST dataset. In “test reconstructions” odd rows correspond to the real test points.



Figure 3: VAE (left column), WAE-MMD (middle column), and WAE-GAN (right column) trained on CelebA dataset. In “test reconstructions” odd rows correspond to the real test points.

Good reconstruction and generation

Gromov-Wasserstein metric [Mémoli (2011)]

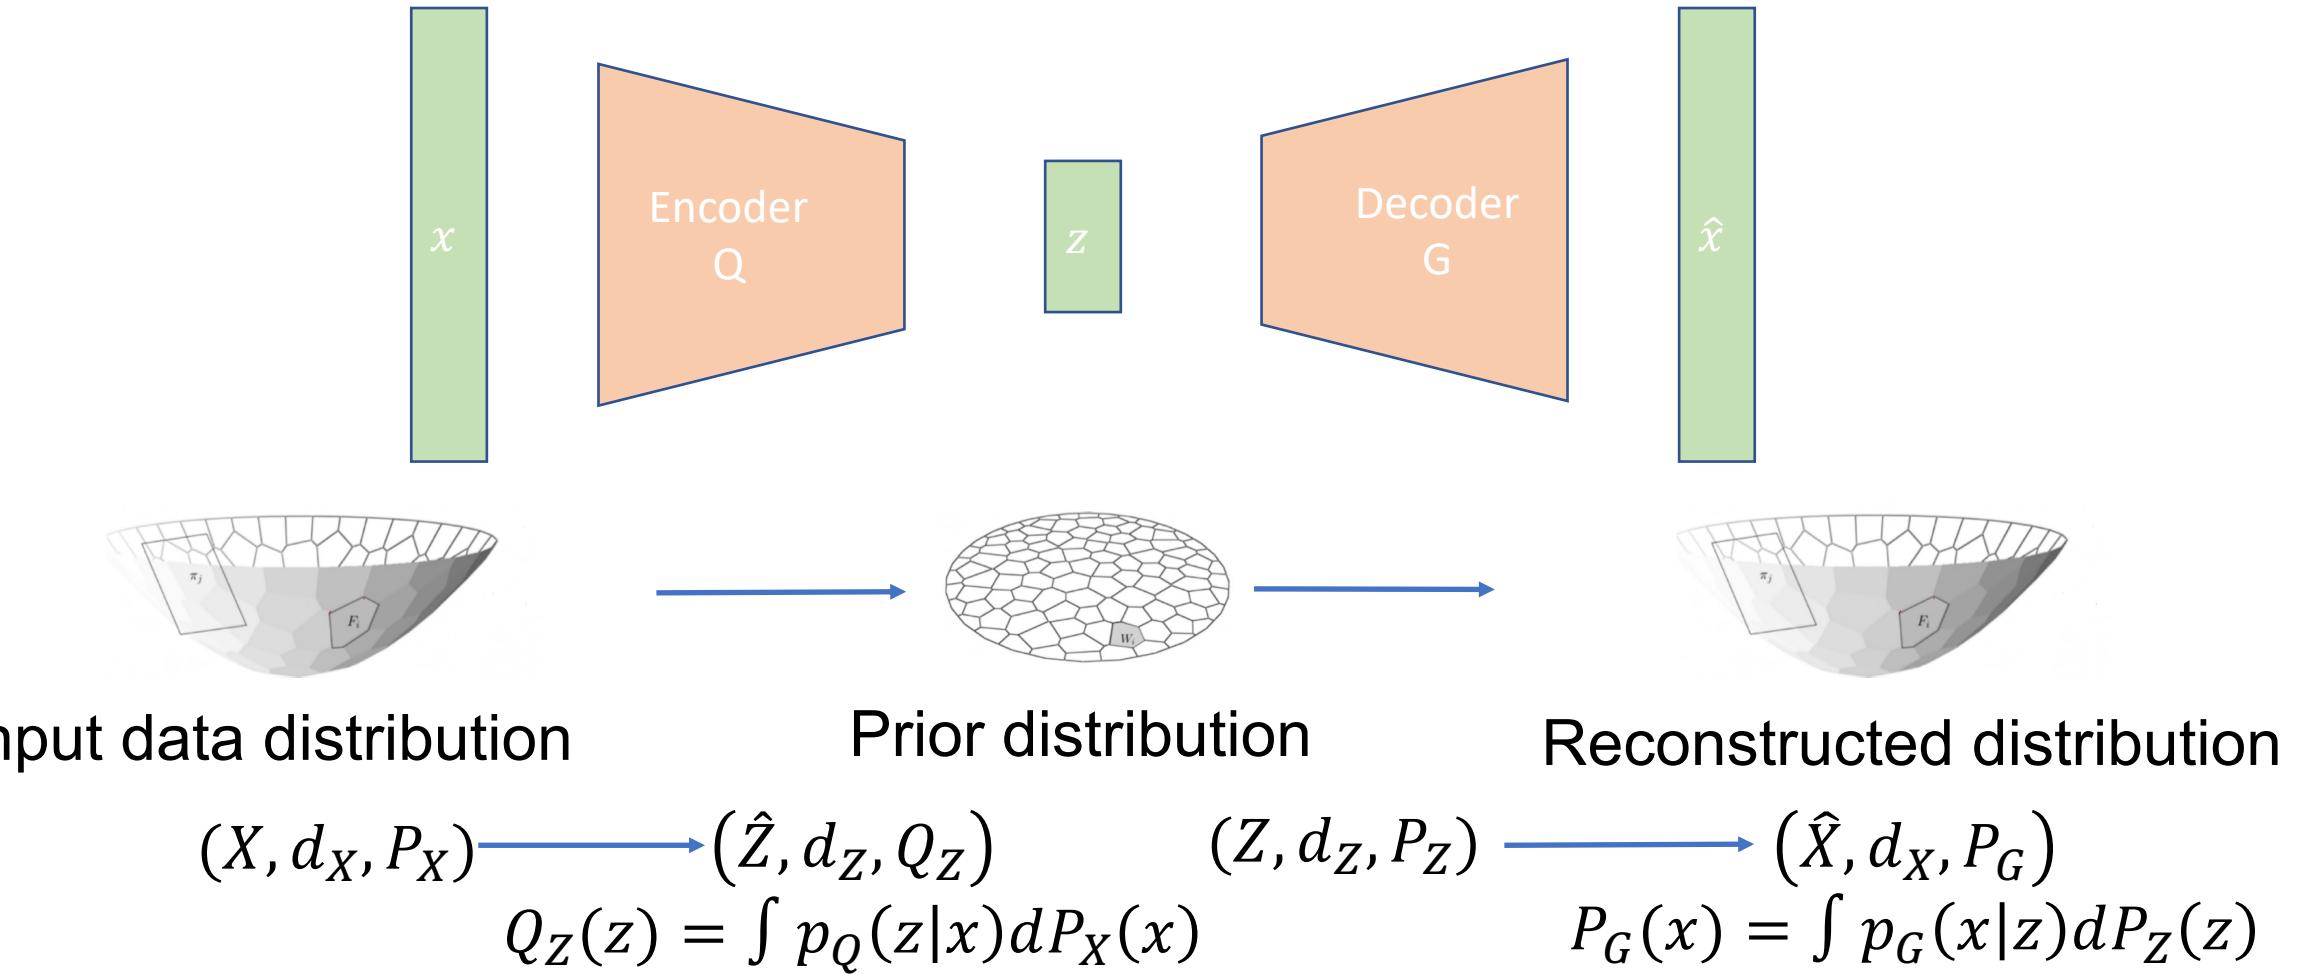
- Wasserstein metric can be regarded as a probabilistic relaxation of the Hausdorff metric:

$$\begin{aligned} d_{\mathcal{H}}^Z(A, B) &:= \max\left(\sup_{a \in A} \inf_{b \in B} d(a, b), \sup_{b \in B} \inf_{a \in A} d(a, b)\right) \\ &= \inf_{R \in \mathcal{R}(A, B)} \sup_{(a, b) \in R} \|a - b\| \\ \mathcal{W}_{\infty}(P, Q) &= \inf_{\mu \in \mathcal{M}(\mu_A, \mu_B)} \sup_{(x, y) \in \text{supp}[\mu]} \|x - y\| \end{aligned}$$

- Gromov-Wasserstein metric is a relaxation of the Gromov-Hausdorff metric:

$$\begin{aligned} d_{\mathcal{GH}}(X, Y) &= \frac{1}{2} \inf_{R \in \mathcal{R}(X, Y)} \sup_{\substack{x_1, x_2 \in X \\ y_1, y_2 \in Y \\ s.t. (x_i, y_i) \in R}} \Gamma_{X, Y}(x_1, y_1, x_2, y_2) \\ d_{\mathcal{GW}, p} &= \frac{1}{2} \inf_{\mu \in \mathcal{M}(\mu_A, \mu_B)} \left(\int \left(\Gamma_{X, Y}(x_1, y_1, x_2, y_2) \right)^p d\mu \otimes d\mu \right)^{1/p} \end{aligned}$$

Gromov-Wasserstein Autoencoder [Gong et al. (2022)]



Minimize $d_{\mathcal{GW},p}$ between original data x and representation z

Gromov-Wasserstein Autoencoders

- Denote the priori distribution by $\pi_\theta(z)$, decoder by $p_\theta(x|z)$ and joint distribution by $p_\theta(x, z) = \pi_\theta(z)p_\theta(x|z)$.
- The Gromov-Wasserstein loss can be written as

$$L_{\mathcal{GW}} := E_{p_\theta(x, z)} E_{p_\theta(x', z')} [\|d_X(x, x') - C d_Z(z, z')\|^k]$$

where C is a scaling parameter.

- To make training stable, **three regularizers** are added:
- Denote the data distribution by $p_{data}(x)$, encoder by $q_\phi(z|x)$ and joint distribution by $q_\phi(x, z) = p_{data}(x)q_\phi(z|x)$. **Wasserstein regularizer** can be written as

$$L_{\mathcal{W}} = E_{q_\phi(x, z)} E_{p_\theta(x'|z)} [d_X(x, x')]$$

Gromov-Wasserstein Autoencoders

- To fulfill the coupling condition in the abovementioned metric, a **sufficient regularizer** is added as:

$$L_D := E_{q_\phi(x,z)}[f_\psi(x,z)] - E_{p_\theta(x,z)}[f_\psi(x,z)]$$

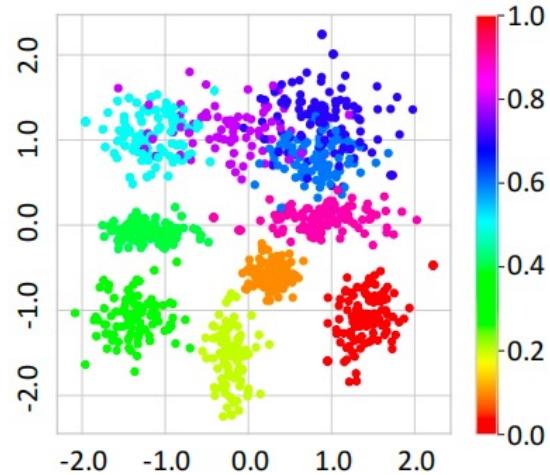
- The **entropy regularization** is used to avoid degenerate solutions

$$R_H := H_q(z|x) = E_{q_\phi(x,z)}[-\log q_\phi(z|x)]$$

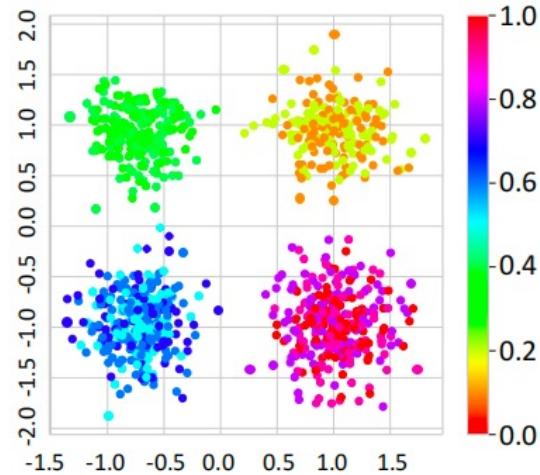
- Then **the total object** can be written as:

$$\min_{\theta, \phi} \max_{\psi} L := L_{\mathcal{GW}} + \lambda_{\mathcal{W}} L_{\mathcal{W}} + \lambda_D L_D - \lambda_H R_H$$

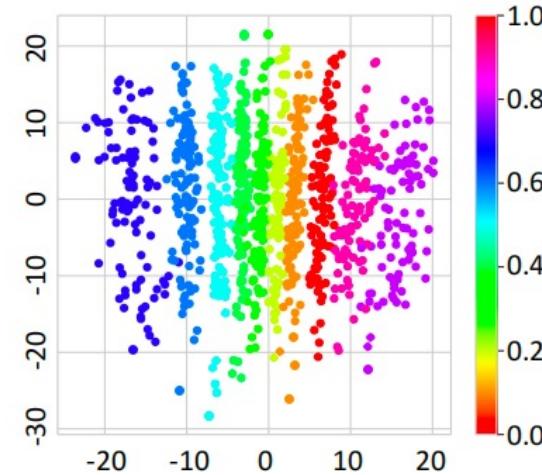
Gromov-Wasserstein Autoencoders



(a) VAE ([Kingma & Welling, 2014](#)).



(b) FactorVAE ([Kim & Mnih, 2018](#)) ($\gamma=10$).



(c) GWAE (FNP, $\lambda_W=1$, $\lambda_D=10$, $\lambda_H=0.3$).

Figure 2: Comparison of the learned latent spaces in 3D Shapes ([Burgess & Kim, 2018](#)) and $L = 16$. The vertical and horizontal axes in the scatter plots respectively represent two of the 16 ($= L$) latent variables with the highest and the second-highest informativeness ([Do & Tran, 2020](#)) w.r.t. the *object hue* factor. Note that a single factor value varies along only one axis in a disentangled representation.

Good representation and clustering

Reference

- 1. Boscaini, D., Masci, J., Rodolà, E., & Bronstein, M. (2016). Learning shape correspondence with anisotropic convolutional neural networks. *Advances in neural information processing systems*, 29.
- 2. Tolstikhin, I., Bousquet, O., Gelly, S., & Schoelkopf, B. (2017). Wasserstein auto-encoders. *arXiv preprint arXiv:1711.01558*.
- 3. Gong, F., Nie, Y., & Xu, H. (2022, October). Gromov-Wasserstein multi-modal alignment and clustering. In *Proceedings of the 31st ACM International Conference on Information & Knowledge Management* (pp. 603-613).
- 4. Mémoli, F. (2011). Gromov–Wasserstein distances and the metric approach to object matching. *Foundations of computational mathematics*, 11, 417-487.
- 5. Ovsjanikov, M., Ben-Chen, M., Solomon, J., Butscher, A., & Guibas, L. (2012). Functional maps: a flexible representation of maps between shapes. *ACM Transactions on Graphics (ToG)*, 31(4), 1-11.
- 6. Bronstein, A. M., Bronstein, M. M., & Kimmel, R. (2006). Generalized multidimensional scaling: a framework for isometry-invariant partial surface matching. *Proceedings of the National Academy of Sciences*, 103(5), 1168-1172.
- 7. Sun, J., Ovsjanikov, M., & Guibas, L. (2009, July). A concise and provably informative multi-scale signature based on heat diffusion. In *Computer graphics forum* (Vol. 28, No. 5, pp. 1383-1392). Oxford, UK: Blackwell Publishing Ltd.

Reference

- 8. Tenenbaum, J. B., Silva, V. D., & Langford, J. C. (2000). A global geometric framework for nonlinear dimensionality reduction. *science*, 290(5500), 2319-2323.
- 9. Qi, C. R., Su, H., Mo, K., & Guibas, L. J. (2017). Pointnet: Deep learning on point sets for 3d classification and segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 652-660).
- 10. Qi, C. R., Yi, L., Su, H., & Guibas, L. J. (2017). Pointnet++: Deep hierarchical feature learning on point sets in a metric space. *Advances in neural information processing systems*, 30.
- 11. Sharma, G., Dash, B., RoyChowdhury, A., Gadelha, M., Loizou, M., Cao, L., ... & Kalogerakis, E. (2022, August). PriFit: learning to fit primitives improves few shot point cloud segmentation. In *Computer Graphics Forum* (Vol. 41, No. 5, pp. 39-50).
- 12. Li, L., Sung, M., Dubrovina, A., Yi, L., & Guibas, L. J. (2019). Supervised fitting of geometric primitives to 3d point clouds. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 2652-2660).
- 13. Litany, O., Remez, T., Rodola, E., Bronstein, A., & Bronstein, M. (2017). Deep functional maps: Structured prediction for dense shape correspondence. In *Proceedings of the IEEE international conference on computer vision* (pp. 5659-5667).
- 14. Halimi, O., Litany, O., Rodola, E., Bronstein, A. M., & Kimmel, R. (2019). Unsupervised learning of dense shape correspondence. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 4370-4379).